


# Programmation Parallèle

**MPI:** Message Passing Interface

---

---

---




# Collective communications

ENSIIE-HPC/BigData-PP-IIP-Lecture 2

---

---

---




# Collectives



- Like p2p, every collective operation requires a communicator
  - Define set of processes that will participate
- All processes in the communicator participate in the collective communication, and have to call the collective function
- All communications which are not p2p are collectives
  - p2p:
    - One-to-one
  - Collectives:
    - One-to-All
    - All-to-One
    - All-to-all

# COLLECTIVE COMMUNICATION



Synchronization

# synchronization




- Communication and synchronization :
  - To ensure the consistency of a calculation, parallel actions can have a « meeting point » before resuming their execution
- These « meeting points » are called **synchronization** :
  - If the synchronization concerns all parallel actions
    - Global or collective synchronization
  - If the actions have different addressing space
    - Communications
- Communications
  - Global synchronization => **global or collective communication**
  - synchronization between 2 actions => **point-to-point communication**

# Barrier


- Synchronize all processes belonging to target communicator

```
int MPI_Barrier( MPI_Comm comm ) ;
```

```
MPI_Init(&argc, &argv);  
  
/* work 1 */  
MPI_Barrier(MPI_COMM_WORLD);  
  
/* work 2 */  
MPI_Finalize();
```




# COLLECTIVE COMMUNICATION



Data Exchange

# Broadcast

- Send data owned by one process to all other processes inside target communicator
- Process emitting data → *root*
- One-to-all collective communication





# Broadcast



```
int MPI_Bcast (  
    void *buf(inout),  
    int count(in),  
    MPI_Datatype datatype(in),  
    int root(in),  
    MPI_Comm comm(in),  
);
```

- rank == **root** → address of memory zone to send
- rank != **root** → address where to store broadcasted data

Output memory segment should be allocated by user.

Size of broadcasted data (number of elements of type **datatype**).

# Broadcast



```
int MPI_Bcast (  
    void *buf(inout),  
    int count(in),  
    MPI_Datatype datatype(in),  
    int root(in),  
    MPI_Comm comm(in),  
);
```

Root rank.

This rank is valid inside **comm** communicator.

All processes involved in this collective should have the same root.

# Broadcast




```
int me, root;
float pi = 0.0 ;
root = 0; /* Process 0 is the root */
...
MPI_Comm_rank(MPI_COMM_WORLD, &me);
...
if (me == root)
    pi = 3.14; /* Only root has the right initial value */
/* All processes have to call MPI_Bcast */
MPI_Bcast(&pi, 1, MPI_FLOAT, root, MPI_COMM_WORLD);

printf("P%d: pi = %f\n", me, pi);
```

```
% srun -n 4 ./a.out
P0: pi = 3.14
P3: pi = 3.14
P1: pi = 3.14
P2: pi = 3.14
%
```

# Scatter

- One root send different data to other processes inside target communicator
- Sent data: same size and same type
- One-to-all collective communication



# Scatter



```
int MPI_Scatter (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcnt(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Start address of memory segment to be sent to each process in **comm** communicator.

**sendbuf** is only valid for root process.

# Scatter



```
int MPI_Scatter (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Message size to send to each target process.

Thus, **sendbuf** array is of size  
**sendcount** × |**comm**| elements.

Elements to be sent to rank  $p$  start at address  
**sendbuf** +  $p \times$  **sendcount**.

# Scatter



```
int MPI_Scatter (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Address where to put received data.

This memory segment should be allocated before calling the collective function.

# Scatter



```
int MPI_Scatter (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Size of message to receive (equals to **sendcount**).

Size is in number of elements of type **rcvtatyp**.



# Scatter



```
int MPI_Scatter (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Root rank.

This rank should be valid inside  
communicator **comm**.

All processes should have the same root.

# Scatter



```
int me, v, root, P;
int *sdbuf;

root = 0; /* Process 0 is the root */

MPI_Comm_rank(MPI_COMM_WORLD, &me);
MPI_Comm_size(MPI_COMM_WORLD, &P);

if (me == root) {
    sdbuf = (int*)malloc(P*sizeof(int));
    fd = fopen("my_data", "r");
    for( p = 0 ; p < P ; p++ )
        sdbuf[p] = read_file_value(fd, p);
    fclose( fd );
} else { sdbuf = NULL; }

MPI_Scatter(sdbuf, 1, MPI_INT,
            &v, 1, MPI_INT, root, MPI_COMM_WORLD);


printf("P%d: received value = %d\n", me, v);
```

```
% cat my_data
18
25
6
3

% srun -n 4 ./a.out
P0: received value = 18
P3: received value = 3
P1: received value = 25
P2: received value = 6
%
```

# Gather

- Root process collects data from other processes (reverse of scatter)
- Sent data: same size and same type
- All-to-one collective communication



# Gather



```
int MPI_Gather (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Address of memory segment  
containing data to send to root  
process.

# Gather



```
int MPI_Gather (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndtotyp(in),  
    void *recvbuf(out),  
    int recvcoun(in)t,  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) ) ;
```

Size (in number of elements of type **sndtotyp**) of message to send to root process.

# Gather



```
int MPI_Gather (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Address of memory segment to receive data from each process inside **comm** communicator.

Only valid for root.

# Gather



```
int MPI_Gather (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Size of data to receive (per process).

Array **recvbuf** must be of size **recvcount** × |**comm**| elements.

# Gather



```
int MPI_Gather (  
    void *sendbuf(in),  
    int sendcount(in),  
    MPI_Datatype sndttyp(in),  
    void *recvbuf(out),  
    int recvcnt(in),  
    MPI_Datatype rcvtatyp(in),  
    int root(in),  
    MPI_Comm comm(in) );
```

Root rank.

This rank is valid inside **comm** communicator

All processes must use the same root.



# Gather



```
int me, v, root, P; int *rcvbuf;  
root = 0; /* Process 0 is the root */
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &me);  
MPI_Comm_size(MPI_COMM_WORLD, &P);
```

```
...  
if (me == root)  
    rcvbuf = (int*)malloc(P*sizeof(int));  
else rcvbuf = NULL;
```

```
v = ... ;  
MPI_Gather(&v, 1, MPI_INT,  
          rcvbuf, 1, MPI_INT, root, MPI_COMM_WORLD);
```


```
if (me == root) {  
    fd = fopen("results", "w");  
    for( p = 0 ; p < P ; p++ )  
        write_file_value(fd, p, rcvbuf[p]);  
    fclose( fd );  
}
```

```
% srun -n 4 a.out  
% cat results  
30  
-100  
23  
19
```

# Allgather

- Equivalent to gather collective operation except that every process involved inside the communication receives the final result.
- Allgather  $\approx$  Gather + Broadcast

|               |       |
|---------------|-------|
| MPI Process 0 | $A_0$ |
| MPI Process 1 | $A_1$ |
| MPI Process 2 | $A_2$ |
| MPI Process 3 | $A_3$ |
| MPI Process 4 | $A_4$ |
| MPI Process 5 | $A_5$ |



|       |       |       |       |       |       |
|-------|-------|-------|-------|-------|-------|
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |
| $A_0$ | $A_1$ | $A_2$ | $A_3$ | $A_4$ | $A_5$ |

# Allgather



```
int MPI_Allgather (
```

```
    void *sendbuf(in),
```

```
    int sendcount(in),
```

```
    MPI_Datatype sndttyp(in),
```

} Arguments corresponding to  
sent data

```
    void *recvbuf(out),
```

```
    int recvcount(in),
```

```
    MPI_Datatype rcvttyp(in),
```


} Arguments corresponding to  
received data

```
    MPI_Comm comm(in),
```

```
);
```

# Alltoall

- Every Process sends and receives data
- Each process send its  $i^{\text{th}}$  pack of data to the  $i^{\text{th}}$  rank
- Realize a sort of matrix transpose



# Alltoall



```
int MPI_Alltoall (
```

```
    void *sendbuf(in),
```

```
    int sendcount(in),
```

```
    MPI_Datatype sendtype(in),
```

```
    void *recvbuf(out),
```

```
    int recvcount(in),
```

```
    MPI_Datatype recvtype(in),
```

```
    MPI_Comm comm(in),
```


```
);
```

} Arguments corresponding to  
sent data

} Arguments corresponding to  
received data

# Reduction

- Root rank collects data from other processes and apply one specific operation
- MPI allows multiple reductions at the same time



# Reduction



```
int MPI_Reduce (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int count(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    int root(in),  
    MPI_Comm comm(in),  
);
```

Address of data to send to root process.

Array of **count** elements of  
**datatype** type.

# Reduction



```
int MPI_Reduce (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int count(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    int root(in),  
    MPI_Comm comm(in),  
);
```

Address of final data (after reduction).

Only valid for root process.

Array of **count** elements of  
**datatype** type.



# Reduction



```
int MPI_Reduce (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int count(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    int root(in),  
    MPI_Comm comm(in),  
);
```

Operation to perform on each element of every process.

Upon return, root will have:

$\text{recvbuf}[i] = \text{op}_{0 \leq p < P} \text{sendbuf}_p[i]$   
with  $0 \leq i < \text{count}$   
and P size of comm.


# Reduction

- Reduction operations are of type `MPI_Op`
- Possibility to create user-defined ones: `MPI_Op_create()`

| <i><b>MPI Operation</b></i> | <i><b>Meaning</b></i> | <i><b>Corresponding Type</b></i> |
|-----------------------------|-----------------------|----------------------------------|
| MPI_MAX                     | Maximum               | Integers and real                |
| MPI_MIN                     | Minimum               | Integers and real                |
| MPI_SUM                     | Sum                   | Integer and real                 |
| MPI_PROD                    | Product               | Integer and real                 |
| MPI_LAND                    | Logical AND           | Integer                          |
| MPI_BAND                    | Binary AND            | Integer and MPI_BYTE             |
| MPI_LOR                     | Logical OR            | Integer                          |
| MPI_BOR                     | Binary OR             | integer MPI_BYTE                 |
| MPI_LXOR                    | Logical XOR           | Integer                          |
| MPI_BXOR                    | Binary XOR            | Integer and MPI_BYTE             |
| MPI_MAXLOC                  | Maximum w/ index      | Structures w/ 2 integers         |
| MPI_MINLOC                  | Minimum w/ index      | Structures w/ 2 integers         |

# Reduce to All

- Equivalent to a reduction
- But every process has the final result
- No need to specify a root process
- Allreduce  $\approx$  Reduce + Broadcast



# Reduce to All



```
int MPI_Allreduce (
```

```
    void *sendbuf(in),
```

```
    void *recvbuf(out),
```

```
    int count(in),
```

```
    MPI_Datatype datatype(in),
```

```
    MPI_Op op(in),
```

```
    MPI_Comm comm(in),
```

```
);
```

Same signature as  
MPI\_Reduce but no root is  
needed.

# Reduce to All

```
int P, N = 100;
int me, i, sum, sum_loc = 0;
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &me);
MPI_Comm_size(MPI_COMM_WORLD, &P);
```

```
...
/* Works if N%P == 0 */
for( i = 1 + me*N/P ; i <= (me+1)*N/P ; i++ )
    sum_loc += i;
```

```
MPI_Allreduce(&sum_loc, &sum_glob, 1, MPI_INT,
              MPI_SUM, MPI_COMM_WORLD);
```


```
/* After reduction, all processes have, in sum_glob, the sum of N first integers */
```

```
printf("1+...+%d = %d\n", N, sum_glob);
```

```
% srun -p 4 ./a.out
1+...+100 = 5050
1+...+100 = 5050
1+...+100 = 5050
1+...+100 = 5050
%
```

# Partial Reduction

- All ranks collect data from other processes **with smaller rank number** and apply one specific operation, including its contribution
- MPI allows multiple reductions at the same time




# Partial Reduction



```
int MPI_Scan (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int count(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    MPI_Comm comm(in),  
);
```

Same signature as  
MPI\_Allreduce

# COLLECTIVE COMMUNICATIONS



Multiple Sizes




# Per-Rank Data Size

- Some collective communications propose multiple versions including one to handle different size for different ranks.
  - E.g., Broadcast, Gather
- Corresponding names have the suffix  $\nabla$ 
  - $\nabla$  = varying or vector or variant
- Examples
  - `MPI_Gather` → `MPI_Gatherv`
  - `MPI_Allgather` → `MPI_Allgatherv`
  - `MPI_Scatter` → `MPI_Scatterv`
  - `MPI_Alltoall` → `MPI_Alltoallv`

# MPI\_Scatterv

- Data  $A_i$  have the same type
- Unlike `MPI_Scatter`,  $A_0, A_1 \dots$  may not have the same size
  - Need another argument to specify size of each data
- Unlike `MPI_Scatter`, data  $A_i$  may not be contiguous (in memory)
  - Need another argument to specify base address of each data



# MPI\_Scatterv

```
int MPI_Scatterv (  
    void *sendbuf(in),  
    int *sendcounts(in),  
    int *displs(in),  
    MPI_Datatype sendtype(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype recvtype(in),  
    int root(in), MPI_Comm comm(in) );
```

**sendbuf** = array of data to distribute to each process inside communicator **comm**.

**sendcounts**[ $p$ ] elements have to send to process  $p$  ( $0 \leq p < P$ ).

Those elements are store at address

**sendbuf** + **displs**[ $p$ ].

All elements are of type **sendtype**.

Those arguments are only valid for root.

# MPI\_Scatterv




```
int MPI_Scatterv (  
    void *sendbuf(in),  
    int *sendcounts(in),  
    int *displs(in),  
    MPI_Datatype sendtype(in),  
    void *recvbuf(out),  
    int recvcount(in),  
    MPI_Datatype recvtype(in),  
    int root(in), MPI_Comm comm(in) );
```

**recvbuf** = array of data to receive.

Array with **recvcount** elements of type  
**recvtype**.

# MPI\_Scatterv



- One process reads a file named data
- This file contains data for each process
- File format is described on the left-hand-side
  - Data size for each process
  - Data (e.g., float values) stored in a per-process manner
- After parsing step, all processes get their own piece of data

# MPI\_Scatterv

```
int sz; int *szbuf, *displs;  
double *sdbuf, *rcvbuf;
```

```
if (me == root) {  
    szbuf = (int*)malloc(P*sizeof(int));  
    displs = (int*)malloc((P+1)*sizeof(int));  
    fd = fopen("data", "r");  
    displs[0] = 0;  
    for( p = 0 ; p < P ; p++ ) {  
        szbuf[p] = read_int_value(fd, p);  
        displs[p+1] = displs[p] + szbuf[p];  
    }
```

*/\* Read data sizes for each  
process (szbuf array)  
Compute corresponding offset  
(displs array) \*/*

```
    sdbuf = (double*)malloc(displs[P]*sizeof(double));  
    for( p = 0 ; p < P ; p++ )  
        for( i = 0 ; i < szbuf[p] ; i++ )  
            sdbuf[displs[p]+i] = read_float_value(fd, p);  
    fclose( fd );  
}
```

*/\* Read per-block  
data (sdbuf array).  
Blocks will be  
distributed among  
ranks \*/*


# MPI\_Scatterv

```
else {
    szbuf = displs = NULL;
    sdbuf = NULL;
}
/* All processes must call MPI_Scatter
   Corresponding data size is stored into sz
   */
MPI_Scatter(szbuf, 1, MPI_INT,
            &sz, 1, MPI_INT,
            root, MPI_COMM_WORLD);

rcvbuf = (double*)malloc(sz*sizeof(double));
/* All process must call MPI_Scatterv
   Corresponding data are stored in rcvbuf
   */
MPI_Scatterv(sdbuf, szbuf, displs, MPI_DOUBLE,
             rcvbuf, sz, MPI_DOUBLE,
             root, MPI_COMM_WORLD);
```


# Reduce-Scatter

- Combine a reduction and a scatter
- Reduction on an array
- Every process gets a piece of array
- Like `MPI_Scatterv`,  $A_0, A_1 \dots$  may not have the same size on the scatter part






# THE THREE MISTAKES



Alltoallw, Exscan, Scatter\_Reduce\_block

# Alltoallw

- Like MPI\_Scatterv,  $A_0, A_1 \dots$  may not have the same size
  - Need another argument to specify size of each data
- Like MPI\_Scatterv, data  $A_i$  may not be contiguous (in memory)
  - Need another argument to specify base address of each data
- In addition, data  $A_i$  may not have the same type
  - Datatype argument is now an array



# MPI\_Alltoallw

```
int MPI_Alltoallw (  
    void *sendbuf(in),  
    int *sendcounts(in),  
    int *displs(in),  
    MPI_Datatype *sendtypes(in),  
    void *recvbuf(out),  
    int *recvcounts(in),  
    MPI_Datatype *recvtypes(in),  
    int root(in), MPI_Comm comm(in) );
```

**sendbuf** = array of data to distribute to each process inside communicator **comm**.

**sendcounts**[*p*] elements have to send to process *p* ( $0 \leq p < P$ ).

Those elements are store at address

**sendbuf** + **displs**[*p*] .

**sendtype**[*p*] is the type of elements to send to process *p*

# MPI\_Alltoallw


```
int MPI_Alltoallw (  
    void *sendbuf(in),  
    int *sendcounts(in),  
    int *displs(in),  
    MPI_Datatype *sendtypes(in),  
    void *recvbuf(out),  
    int *recvcounts(in),  
    MPI_Datatype *recvtypes(in),  
    int root(in), MPI_Comm comm(in) );
```

**recvbuf** = array of data to receive.

Array composed of **recvcount**[*p*] elements  
of type **recvtype**[*p*] coming from all other  
processes *p*.

# Reduce-Scatter block

- Combine a reduction and a scatter
- Reduction on an array
- Every process gets a piece of array of same size
- Data chunk are contiguous (no displacement arg)



# Reduction-Scatter block




```
int MPI_Reduce_Scatter_block (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int recvcnt(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    MPI_Comm comm(in),  
);
```

Address of data to send to root process.

Fixed **recvcnt** elements of  
**datatype** type.

# Partial Reduction

- All ranks collect data from other processes with smaller rank number and apply one specific operation, excluding its contribution
- MPI allows multiple reductions at the same time



# Partial Reduction



```
int MPI_Exscan (  
    void *sendbuf(in),  
    void *recvbuf(out),  
    int count(in),  
    MPI_Datatype datatype(in),  
    MPI_Op op(in),  
    MPI_Comm comm(in),  
);
```


Same signature as  
MPI\_Allreduce



# Why mistakes

- Reduce\_scatter\_block is the simplified version of reduce\_scatter
  - should have been the first one to appear
  - For coherency in naming convention, reduce\_scatter\_block should be reduce\_scatter, and reduce\_scatter should be reduce\_scatterv
- MPI\_Scan is just (MPI\_Exscan + local variable)
  - Do not need to create two functions
  - Function MPI\_Exscan is enough. Should have been the first one
- MPI\_Alltoallw
  - Allows the user to exchange anything
  - Multiple datatypes -> complex internal implementation and algorithms
  - Why is it the only collective to allow multiple datatypes?

# COLLECTIVES



Usage

# Collectives



- Completing a collective call on a process does not mean that the collective is done for all processes
  - It just means that the contribution of the current process is done
- If the current process does not require the result but only send data, it basically is the same as MPI\_Send
  - After the collective call, the user can reuse the input buffer
  - It does not even mean that the data has been send
- If all processes provide input data and require an output result, the collective is synchronizing
  - However, Barrier is the only official synchronizing collective in the API

# Collectives



- Every collective operation requires a communicator
  - Define set of processes that will participate
- What happen if two ranks call a different collective operation?
- What happen if two ranks in different communicator call the same collective communication at the same time?

# Rules

- 
- 
- All processes in the group identified by the communicator must call the collective routine
  - Inside a communicator, processes should call the same sequence of collective communication
  - Between two ranks in different communicators, there are no restrictions

# Examples

```
void f ( int r ) {  
    if( r == 0 )  
        MPI_Barrier(MPI_COMM_WORLD);  
    return;  
}
```

May be wrong  
(depends on  
value of r)

```
void g ( int r ) {  
    if( r == 0 )  
        MPI_Barrier(MPI_COMM_WORLD);  
    else  
        MPI_Barrier(MPI_COMM_WORLD);  
    return;  
}
```

Correct


```
void h ( int r ) {  
    if( r == 0 ) {  
        MPI_Reduce(MPI_COMM_WORLD, ...);  
        MPI_Barrier(MPI_COMM_WORLD);  
    } else {  
        MPI_Barrier(MPI_COMM_WORLD);  
        MPI_Reduce(MPI_COMM_WORLD, ...);  
    }  
    return; }
```

May be wrong  
(depends on  
value of r)

# Hints

- 
- Be careful of arguments to MPI collective communications
    - Same root
    - Compatible count for elements to send and to receive
  - Be careful of control flow inside the program
    - Not related to lexical order
    - Related to dynamic order
  - Be careful of communicators


# COLLECTIVE ALGORITHMS






# Broadcast algorithms

- Linear algorithm
- Root send its data to each other process
- Iterative and linear: for  $p$  processes, need  $p-1$  step




# Broadcast algorithms

- Binary tree: at each step, each process with the data to broadcast send the data to another process
- Need  $\log(p)$  steps
- For(0;  $i < \log(p)$ ;  $i++$ ) if(rank  $\leq i$ ) send(data, rank +  $2^i$ )




# Broadcast algorithms

- Binary tree
- For( $\log(p)$ ;  $i > 0$ ;  $i--$ ) if( $\text{rank} \% 2^i == 0$ ) send(data,  $\text{rank} + 2^{i-1}$ )
- Are those two algorithms equivalent?




# Broadcast algorithms

- If we consider two nodes



# Broadcast algorithms

- If we consider two nodes




# Broadcast algorithms

- If we consider two nodes
- All internode communications ( network) happens at the same step

**Node 1**

**Node 2**



# Broadcast algorithms

- Only one internode communication
- No contention on the network
- Intranode comm may be optimized

**Node 1**


**Node 2**

1 internode comm

No internode comm

No internode comm

No internode comm



# Collective algorithms



- Several algorithms may be implemented for one collective in a runtime
- Default selection of algorithms often depends on the size of the messages, and the number of MPI processes involved
- Algorithm selection may also be available to user through launch options, environment variables or communicator attributes